MOONEY FACE CLASSIFICATION AND PREDICTION BY LEARNING ACROSS TONE

Tsung-Wei Ke⋆†   *Stella X. Yu*⋆†   *David Whitney*⋆

⋆ UC Berkeley / †ICSI

ABSTRACT

Mooney faces are special two-tone images that elicit a rich impression of identity and facial expression in human observers. While Mooney faces are important, there exist only a small number of instances hand-crafted from source photos which are often no longer available.

We first apply deep learning methods to generate a plausible Mooney face automatically from any face photo. We are then able to create a large-scale face dataset with paired grayscale and two-tone images.

We then study how well two-tone versions make face predictions, using conditional Generative Adversarial Networks. We show that faces predicted from Mooney images bear striking resemblance to source photos, and they are better than two-tone images obtained by global intensity thresholding. We also demonstrate remarkable face predictions from very low resolution surveillance photos. Our findings reveal great potentials of combining deep learning and Mooney faces for more effective face recognition in a wide range of conditions.

Index Terms— Mooney Faces, Generative Adversarial Network, Cross-Tone Integration

1. INTRODUCTION

Mooney faces are special two-tone black and white pictures of faces (Fig.1). Despite the lack of detail, shading, and other cues, human observers can often effortlessly recognize the identity, age, gender, and facial expression of Mooney faces [1, 2, 3]. Even infants can recognize Mooney faces [4, 5, 6]. How such a remarkable perceptual ability is achieved with such sparse information may be key to understanding human face recognition in naturalistic settings, where noise, occlusion, and shadows are common.

Mooney faces may be challenging for both human and computer vision systems because the contours are ambiguous; it is not clear how they can be segmented into parts based on bottom-up image-level information alone. While it is trivial to extract the contours of a shadowy Mooney face, it is not obvious how to interpret or classify these contours. In a two-tone Mooney face, each contour could correspond to a meaningful structure of the face (what is sometimes called an *attached shadow*) or to a less meaningful *cast shadow* border. Cast shadow borders depend on the nature and direction of the light

Fig. 1. Mooney faces are special two-tone images where humans can often effortlessly recognize the identity, age, gender, and facial expression of the person. It remains elusive how such rich perception can be achieved computationally.

source, the shape of the intervening surface, and the surface on which the shadow is cast. Thus, there is skepticism about the amount of usable information from cast shadow borders for human recognition [7]. Attached shadows depend primarily on the structure of the surface, and are therefore useful for recognition. However, parsing the contour into relevant attached shadow segments and irrelevant cast shadow segments requires some prior knowledge about the image. For these reasons, it is thought that Mooney faces require holistic [2] or top-down information for recognition [2, 3], for example by template matching [8].

Face recognition from photos is well studied in computer vision [9, 10], however, current approaches applied out of the box would fail miserably on two-tone Mooney faces. If there were an algorithm for recognizing Mooney faces, it could shed light on human visual processing. It could also shed light on broader issues in computer vision like face recognition in noise, occlusion, and extreme lighting conditions, which are frequently encountered in surveillance videos.

Mooney faces are not well defined; not every binarized face image appears as a Mooney face. Mooney faces are usually constructed by artists or scientists manually, in an ad-hoc manner, requiring selective grayscale image editing and human subjective judgment, a time and labor-consuming process. Such a mechanism for generating Mooney faces does not scale, severely limiting current approaches of studying Mooney faces. Our first goal here is to generate a best possible Mooney face automatically from a face photo. We take a

deep learning approach to train a Mooney face classifier using a limited number of Mooney faces but many more grayscale face and non-face photos.

Once we have the ability to generate a plausible Mooney face from a photo, we essentially create paired data between grayscale face photos and their Mooney faces. This is an advantage over traditional Mooney faces, as the source photos for previously existing Mooney faces are no longer available. Our second goal here is to use our large-scale paired photo-Mooney data to train a patch-based conditional Generative Adversarial Network (GAN) model for recovering grayscale face photos from two-tone Mooney images. The grayscale faces predicted by our model on novel Mooney faces are strikingly similar to their source images, even though not a single photo of the same person has been used during training.

We next study how special Mooney faces are among two-tone images. We compare faces predicted from Mooney images and from those two-tone images obtained simply by global intensity thresholding on source photos. We observe that Mooney-to-Photo predictions are better, indicating Mooney images may carve out a special image space that retains the minimal stable structures of faces.

We apply our Photo-to-Mooney generation Mooney-to-Photo prediction models successively in order to predict faces from very low resolution surveillance photos. Our remarkable results reveal the great potentials of combining Mooney faces and deep learning for more robust and effective face recognition.

2. MOONEY FACE CLASSIFIER AND DATASET

We develop a deep learning model to automatically generate a plausible Mooney face from a given face photo (Fig.2). We create two-tone Mooney candidates by smoothing and thresholding the face photo with multiple parameter settings. We train a face classifier to determine whether a two-tone image is a Mooney face. We feed the Mooney face candidates into our classifier and the top scorer becomes our automatically generated Mooney face.

The key to our approach is how to train our Mooney face classifier, since we only have a very limited number of Mooney faces found in the literature. We take a two-step approach. We first use large scale FaceScrub dataset [11] and non-face ImageNet dataset [12] to pre-train a grayscale face classifier. We then fine-tune the model with the small number of Mooney faces using various data augmentation techniques.

Grayscale face classifier. We build our collection of positive and negative face image examples to train a grayscale face classifier. Our face photos are 90,000 FaceScrub images of about 500 identities, whereas our non-face photos are 90,000 images uniformly sampled over 1,000 classes of the ImageNet dataset. We extract the model from the triplet network of OpenFace [13] and append it with a two-way softmax layer for representing the output of whether the input image


Fig. 2. Overview of our automatic Mooney face generator. Given a grayscale image, we generate a set of black/white versions by smoothing and thresholding, each of which is evaluated by our Mooney classifier and the highest scorer is the final Mooney face.

is a face or not. The input image resolution is set to 96×96. We optimize the network with stochastic gradient descent for 40,000 iterations, with mini-batch size 128, momentum 0.9, and weight decay 0.0005. The learning rate starts at 0.01 and exponentially decays to 0.001.

Mooney face classifier. There are only a few Mooney faces publicly available. Our Mooney face examples come from 48 images of the PICS dataset [14] and our own 80 Mooney faces. It is also known in visual psychology that Mooney faces are not recognizable when the intensity polarity is reversed or the image is inverted. We compose our non-Mooney examples of five types of images: 90,000 binarized ImageNet images, their negative versions (where black pixels become white pixels, and vice versa), the Mooney face inversions, the Mooney face negatives, and their inverted versions.

We fine-tune the previously trained grayscale face classifier for Mooney-face classification. We follow the same hyper-parameter setting for training and testing. We also perform additional data processing for augmentation: Mooney faces are randomly dilated or eroded; ImageNet images are augmented with Photo-to-Mooney data processing.

Photo-to-Mooney data processing. We generate Mooney candidates from a source photo by first smoothing it with a set of $k \times k$ kernels by convolution and then binarizing it with thresholds t, where $k \in \{2, 3, 4, 5, 6\}$ and $t \in \{0.4, 0.45, 0.475, 0.5, 0.525, 0.55, 0.6\}$. The pixel intensity is rescaled to the range between 0.4 and 0.6 before thresholding. There could be other alternatives; the goal is simply to generate relatively fewer black-white blobs that tend to be more Mooney-like.

Fig. 3. Sample results from our Mooney face generator.

Fig. 4. Sample Mooney-to-Photo prediction results.

Method	L_2 error per pixel: mean $\pm$ std
Baseline CCA-256	0.0399 ± 0.0003
Baseline CCA-3200	0.0446 ± 0.0003
Baseline PCCA-256	0.0173 ± 0.0005
Our Patch CGAN	$\mathbf{0.0156 \pm 0.0006}$

Table 1. Mooney-to-face prediction errors on the test set over intensity range 1. For CCA-k, k is the number of CCA bases. PCCA denotes patch-wise CCA, patch size 16×16.

Large-scale Mooney face dataset. We generate 35 candidates for each source face photo from the 35 combinations of k and t. Our Mooney face classifier is used to pick out the most Mooney-like candidate. If that candidate has a low score, we reject the face photo. We build our large-scale Mooney face dataset based on Facescrub images. With a high Mooney score threshold to ensure the quality of such automatically generated Mooney faces without human judgement verification, we obtain a Mooney face dataset with 13,460 images of 523 identities, sample results shown in Fig 3.

3. MOONEY-TO-PHOTO PREDICTION

Our large-scale Mooney face dataset with paired source photos and Mooney faces opens up exciting new opportunities. We explore predicting fine-tone faces from two-tone Mooney faces, emulating human vision's rich face perception from two-tone Mooney faces and enabling computer face recognition in low-bit images from a wide range of conditions.

Related Works. Mooney-to-photo prediction is an interesting and challenging research topic in computer vision,

and only a few works have attempted solving the problem. Shashua *et al.*[15] use linear models to recover faces from two-tone images under several restrictive assumptions, e.g. Lambertian reflection models and 3D face models. Maver *et al.*[16] build the eigenspace from a set of grayscale faces, use linear models to predict the coefficients from input two-tone images and then reconstruct the grayscale faces. Kemelmacher *et al.*[17] propose to reconstruct faces from Mooney images with prior knowledge of 3D face models. We approach the problem from a pure data-driven deep learning perspective, translating a Mooney to a face image directly.

Our Patch CGAN Model. Generative Adversarial Networks (GAN) *et al.*[18] consist of a Discriminator (D) and a Generator (G). G is trained to produce fake data that looks real, whereas D learns to classify data as real or fake. This idea is extended to so-called Conditional GAN (CGAN), where both D and G are conditioned on extra data [19].

Isola *et al.*applies CGAN to paired input data for cross-modal translation from one type of images to another [20]. Their models use an encoder-decoder net with skip connections for G, which first encodes an input image and then de-

Fig. 5. Sample prediction results from two-tone inputs.

	two-tone inputs	L_2 error per pixel: mean $\pm$ std
1-bit	30%-white	0.0199 ± 0.0005
1-bit	40%-white	0.0189 ± 0.0005
1-bit	50%-white	0.0175 ± 0.0005
1-bit	Mooney	$\mathbf{0.0156 \pm 0.0006}$
3 1-bit	(30,40,50)%-white	0.0148 ± 0.0005

Table 2. Two-tone to grayscale face prediction errors on the test set over the intensity range of 1.

codes the latent representation to generate an output image. They use a convolutional Patch CGAN classifier for D, collecting real/fake scores over 16×16 patches. We follow the same Patch CGAN setting to train a CGAN on paired (Mooney,photo) data in order to predict a plausible grayscale face photo from a Mooney image.

Baseline CCA methods. Since our D is based on 16×16 patches and our G is based on 128×128 images, we apply correlation coefficient analysis (CCA) at both batch and full image levels as linear baseline models.

Experimental Setup. We split our new Mooney face dataset evenly into train, validation and test sets by identity. We resize the shorter size of each input image to 128 and randomly crop a 128×128 patch for training, whereas we only use center crops for testing. Horizontal flipping is used for additional data augmentation. We use ADAM [21] to optimize the network. We set the initial learning rate to 0.0002, momentum 0.5, batch size 1, λ_{L1} 100.

Cross-Tone Prediction Results. Fig.4 shows grayscale face predictions and Table3 prediction errors. Our CGAN

Fig. 6. Sample prediction results from surveillance photos.

model significantly outperforms CCA models. Our CGAN predictions are strikingly similar to source photos, despite the fact that our network has never seen any faces of the test identities and the inputs are merely two-tone Mooney images. The linear CCA models have trouble generalizing to unseen faces: full-image CCA produces grainy results, whereas patch-wise CCA produces over-smoothed results which lack highlights and shadows and look more like two-tone Mooney faces.

Are Mooney Faces Special? We generate alternative two-tone images by intensity thresholding. We binarize each photo by an intensity percentile. Since ground-truth Mooney faces have a mean white pixel percentage of 40%, we choose three percentiles: 30%, 40%, 50%. We also concatenate these binary versions into a color (3 1-bit) image. We follow the same setting to train and test face predictions from such two-tone images. Fig. 5 shows prediction results and Table3 prediction errors. Mooney faces produce better and non-trivial predictions than simple two-tone images.

Face Prediction from Very Low-Resolution Surveillance Photos. We crop 12×12 face patches from the VIPeR surveillance dataset[22] and resize them to 128×128. We apply our FacePhoto-to-Mooney model to generate its Mooney version, and then apply our Mooney-to-Photo model to predict a grayscale photo. Fig 6 shows remarkable 128×128 faces predicted from 12×12 face patches.

Summary. We develop deep learning models for Photo-to-Mooney generation and Mooney-to-face prediction. We build the first large-scale paired (face,Mooney) dataset, and learn two-tone Mooney to full-tone face predictions. We show that Mooney images produce better and non-trivial face predictions than intensity thresholded two-tone images. We apply both our models successively to predict faces from extreme low-resolution face images. These results are useful for understanding what makes Mooney faces special and for developing better face recognition methods.

4. REFERENCES

[1] Craig M Mooney, "Age in the development of closure ability in children.," *Canadian Journal of Psychology/Revue canadienne de psychologie*, vol. 11, no. 4, pp. 219, 1957.

[2] Daphne Maurer, Richard Le Grand, and Catherine J Mondloch, "The many faces of configural processing," *Trends in cognitive sciences*, vol. 6, no. 6, pp. 255–260, 2002.

[3] Elinor McKone, "Isolating the special component of face recognition: peripheral identification and a mooney face.," *Journal of Experimental Psychology: Learning, Memory, and Cognition*, vol. 30, no. 1, pp. 181, 2004.

[4] Faraz Farzin, Susan M Rivera, and David Whitney, "Spatial resolution of conscious visual perception in infants," *Psychological science*, vol. 21, no. 10, pp. 1502–1509, 2010.

[5] Irene Leo and Francesca Simion, "Newborns' mooney-face perception," *Infancy*, vol. 14, no. 6, pp. 641–653, 2009.

[6] Yumiko Otsuka, Harold CH Hill, So Kanazawa, Masami K Yamaguchi, and Branka Spehar, "Perception of mooney faces by young infants: The role of local feature visibility, contrast polarity, and motion," *Journal of Experimental Child Psychology*, vol. 111, no. 2, pp. 164–179, 2012.

[7] Cassandra Moore and Patrick Cavanagh, "Recovery of 3d volume from 2-tone images of novel objects," *Cognition*, vol. 67, no. 1, pp. 45–71, 1998.

[8] Patrick Cavanagh, "Whats up in top-down processing," *Representations of vision: Trends and tacit assumptions in vision research*, pp. 295–304, 1991.

[9] Florian Schroff, Dmitry Kalenichenko, and James Philbin, "Facenet: A unified embedding for face recognition and clustering," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2015, pp. 815–823.

[10] Yandong Wen, Kaipeng Zhang, Zhifeng Li, and Yu Qiao, "A discriminative feature learning approach for deep face recognition," in *European Conference on Computer Vision*. Springer, 2016, pp. 499–515.

[11] Hong-Wei Ng and Stefan Winkler, "A data-driven approach to cleaning large face datasets," in *Image Processing (ICIP), 2014 IEEE International Conference on*. IEEE, 2014, pp. 343–347.

[12] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei, "ImageNet Large Scale Visual Recognition Challenge," *International Journal of Computer Vision (IJCV)*, vol. 115, no. 3, pp. 211–252, 2015.

[13] Brandon Amos, Bartosz Ludwiczuk, and Mahadev Satyanarayanan, "Openface: A general-purpose face recognition library with mobile applications," Tech. Rep., CMU-CS-16-118, CMU School of Computer Science, 2016.

[14] Psychological Image Collection at Stirling, "2d face sets," http://pics.stir.ac.uk.

[15] Amnon Shashua, "On photometric issues in 3d visual recognition from a single 2d image," *International Journal of Computer Vision*, vol. 21, no. 1-2, pp. 99–122, 1997.

[16] Jasna Maver and Aleš Leonardis, "Recognizing 2-tone images in grey-level parametric eigenspaces," *Pattern recognition letters*, vol. 23, no. 14, pp. 1631–1640, 2002.

[17] Ira Kemelmacher-Shlizerman, Ronen Basri, and Boaz Nadler, "3d shape reconstruction of mooney faces," in *Computer Vision and Pattern Recognition, 2008. CVPR 2008. IEEE Conference on*. IEEE, 2008, pp. 1–8.

[18] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio, "Generative adversarial nets," in *Advances in Neural Information Processing Systems 27*, Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger, Eds., pp. 2672–2680. Curran Associates, Inc., 2014.

[19] Mehdi Mirza and Simon Osindero, "Conditional generative adversarial nets," *arXiv preprint arXiv:1411.1784*, 2014.

[20] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A Efros, "Image-to-image translation with conditional adversarial networks," *arxiv*, 2016.

[21] Diederik Kingma and Jimmy Ba, "Adam: A method for stochastic optimization," *arXiv preprint arXiv:1412.6980*, 2014.

[22] Douglas Gray and Hai Tao, "Viewpoint invariant pedestrian recognition with an ensemble of localized features," *Computer Vision–ECCV 2008*, pp. 262–275, 2008.